



**GPU Architectures
and Programming
Assignment- Week3**

TYPE OF QUESTION: MCQ/MSQ

**Number of questions: 10
10**

Total mark: 10 X 1 =

MCQ Question

Question 1:

Memory management function for allocating memory in GPU device using CUDA is:

- A. malloc
- B. cudaMemcpy
- C. cudaMalloc
- D. calloc

Ans: B

Question 2:

Correct statement for allocation of memory chunk for 1024 floating point data element in GPU device is (Assume mem_chunk is device pointer):

- A. cudaMalloc((void **)&mem_chunk, 1024*sizeof(float))
- B. cudaMalloc((void **)&mem_chunk, 1024)
- C. cudaMalloc((float **)&mem_chunk, 1024*sizeof(float))
- D. cudaMalloc((float **)&mem_chunk, 1024)

Ans: A

Question 3:

Which of the following CUDA kernels correctly multiplies two 1D arrays element-wise:

```
A. __global__ void multiplyArrays(float *A, float *B, float *C, int n) {  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    if (i < n) {  
        C[i] = A[i] * B[i];  
    }  
}
```

```
B. __global__ void multiplyArrays(float *A, float *B, float *C, int n) {  
    int i = threadIdx.x;  
    C[i] = A[i] * B[i];  
}
```

```
C. __global__ void multiplyArrays(float *A, float *B, float *C, int n) {  
    int i = blockIdx.x * blockDim.x;  
    if (i < n) {  
        C[i] = A[i] * B[i];  
    }  
}
```

```
D. __global__ void multiplyArrays(float *A, float *B, float *C, int n) {  
    int i = threadIdx.x + threadIdx.y * blockDim.x;  
    if(i < n)  
        C[i] = A[i] * B[i];  
}
```

Ans: A

Question 4:

What is the output of the following code if a kernel fails during execution:

```
cudaError_t err = cudaGetLastError();  
if (err != cudaSuccess) {  
    printf("CUDA error: %s\n", cudaGetErrorString(err));  
}
```

- A. It prints the error code as an integer.
- B. It prints the error message corresponding to the error code.
- C. It resets the last error state and continues execution.
- D. It terminates the program immediately.

Ans: B

Question 5:

Which among the following statements is True:

- A. `__device__` function can have return type other than void
- B. `__global__` function must always returns int
- C. A function can not be declared as both `__host__` and `__device__` function
- D. Every function is a default `__global__` function

Ans: A

Question 6:

How is the total number of blocks to be launched determined in the code:

- A. $n / \text{threadsPerBlock}$
- B. $\text{threadsPerBlock} / n$
- C. $(n + \text{threadsPerBlock} - 1) / \text{threadsPerBlock}$
- D. $(n - \text{threadsPerBlock}) / \text{threadsPerBlock}$

Ans: C

Question 7:

Consider a vector addition kernel launch
`vectorAdd<<<dim3(16, 8), dim3(32, 16)>>>(d_A, d_B, d_C, n);`

What does the following CUDA kernel launch configuration imply:

- A. 128 threads per block and 128 blocks in the grid.
- B. 256 threads per block and 128 blocks in the grid.
- C. 512 threads per block and 2048 blocks in the grid.
- D. 512 threads per block and 128 blocks in the grid.

Ans: D

Solution:

Total number of blocks in the grid : $16 * 8 = 128$

Total number of threads in each block: $32 * 16 = 512$

Question 8:

A typical CUDA program structure consists of five main steps:

1. Allocate GPU memories.
2. Copy data from CPU memory to GPU memory.
3. Invoke the CUDA kernel to perform program-specific computation.
4. Copy data back from GPU memory to CPU memory.
5. Destroy GPU memories.

CUDA API for the step no 1,2 and 5 respectively are

- A. `cudaMalloc()`, `cudaMemcpy()`, `cudaFree()`
- B. `malloc()`, `copy()`, `free()`
- C. `cudaAlloc()`, `cudaCopy()`, `cudaRelease()`
- D. `cudaAlloc()`, `copy()`, `free()`

Ans: A

Question 9:

Consider the following kernel specific system variables and the subsequent statements (a)-(d).

- (i) `gridDim.x`
- (ii) `blockDim.y`
- (iii) `blockIdx.y`
- (iv) `threadIdx.x`

- (a) number of blocks in dimension x of multi-dim grid
- (b) thread number inside a block in dimension x
- (c) number of threads per block in dimension y of multi-dim block.
- (d) block number for a thread in dimension y

Match and pair. Choose the correct option that represents the correct solution.

- A. (i)-(a), (ii)-(c), (iii)-(d), (iv)-(b)
- B. (i)-(a), (ii)-(b), (iii)-(c), (iv)-(d)
- C. (i)-(d), (ii)-(a), (iii)-(b), (iv)-(c)
- D. (i)-(a), (ii)-(c), (iii)-(b), (iv)-(d)

Answer A

Question 10:

A CUDA kernel is launched with the following configuration:
`vectorMultiply<<<dim3(8, 2, 4), dim3(16, 32, 2)>>>(d_A, d_B, d_C, n);`
How many total threads will be launched:

- A. 4096
- B. 16,384
- C. 32,768
- D. 65,536

Ans: D

Solution:

Total number of threads launched: $8 \cdot 2 \cdot 4 \cdot 16 \cdot 32 \cdot 2 = 65,536$

*******END*******