



**GPU Architectures
and Programming
Assignment- Week 4
TYPE OF QUESTION: MCQ/MSQ**

Number of questions: 10

Total mark: 10 X 1 = 10

MCQ / MSQ Question

Question 1:

Consider a 1D CUDA kernel that computes on a 1D array A of size 2048 with launch parameters $\langle\langle\langle 32 \rangle, 64 \rangle\rangle\rangle$. Each CUDA thread operates only on one data point. On what data point will the thread with $\text{threadIdx.x}=20$, $\text{blockIdx.x}=20$ operate on. Assume that thread ids along the x-dimension represent rows and thread ids along the y-dimension represent columns. Choose the correct option.

- A. A[200]
- B. A[40]
- C. A[1300]
- D. None of the above

Ans: C

Solution:

Grid configuration: $\langle\langle\langle 32, 1, 1 \rangle\rangle\rangle$

This means there are 32 blocks along x direction

Block configuration: $\langle\langle\langle 64, 1, 1 \rangle\rangle\rangle$

This means there are 64 threads along the x-dimension(1D)

Total threads = $64 \times 32 = 2048$, this matches the size of A

Global thread Id = $\text{threadIdx.x} + \text{blockIdx.x} * \text{blockDim.x}$

Global thread Id = $20 + 64 * 20 = 20 + 1280 = 1300$

So it will operate on data point = A[1300]

Question 2:

Consider a kernel launched with the following launch configuration parameters.

<<<(2,2,1),(32,32,4)>>>

What is the total number of CUDA threads that are spawned for the given launch configuration?

- A. 16384
- B. 2048
- C. 8192
- D. 4096

Ans: A

Solution:

Total number of threads: $2*2*1*32*32*4 = 16384$

Question 3:

For the launch parameter configuration provided in Question 2, what are the minimum and maximum values of threadIdx.x?

- A. min=0, max=1023
- B. min=1,max=32
- C. min=0,max=31
- D. min=1,max=1024

Answer: C

Solution:

We can see that blockDim.x = 32 i.e the number of threads in in x dimension of block is 32. So id of the threads in the block will range from 0 to 31.

Question 4:

For the launch parameter configuration provided in Question 2, what are the minimum and maximum values of blockIdx.y?

- A. min=0,max=1
- B. min=1,max=2
- C. min=0,max=31
- D. min=1,max=32

Answer: A

Solution:

We can see that $\text{blockDim.y} = 32$ i.e the number of threads in in y dimension of block is 32. So id of the threads in the block will range from 0 to 31.

Question 5:

A CUDA kernel is launched with a following grid and block configuration. Given:

Grid configuration: $\langle\langle\langle 2, ?, 1 \rangle\rangle\rangle$

Block configuration: $\langle\langle\langle 4, 4, ? \rangle\rangle\rangle$

With what values should ? be replaced if the total number of threads are 128:

- A. Grid: $\langle\langle\langle 2, 4, 1 \rangle\rangle\rangle$, Block: $\langle 4, 4, 2 \rangle$
- B. Grid: $\langle\langle\langle 2, 4, 1 \rangle\rangle\rangle$, Block: $\langle 4, 4, 1 \rangle$
- C. Grid: $\langle\langle\langle 2, 2, 1 \rangle\rangle\rangle$, Block: $\langle 4, 4, 2 \rangle$
- D. Grid: $\langle\langle\langle 2, 8, 1 \rangle\rangle\rangle$, Block: $\langle 4, 4, 2 \rangle$

Ans: B

Solution:

Total number of threads $2*4*4*4 = 128$

Question 6:

What will happen if `__syncthreads()` is called conditionally, as shown in the code snippet :

```
__global__ void conditionalSync() {  
    if (threadIdx.x % 2 == 0) {  
        __syncthreads();  
    }  
}
```

- A. All threads synchronize correctly.
- B. Only even threads synchronize correctly.
- C. Kernel execution may lead to undefined behavior.
- D. No synchronization takes place.

Ans: C, D

Solution:

__syncthreads() function must be called by all the threads in the block. But here threads with only even id calls the synchronization function. This leads to undefined behavior.

Question 7:

Consider the given kernel code:

```
__global__ void syncExample(int *array) {  
    int idx = threadIdx.x + blockIdx.x * blockDim.x;  
    if (threadIdx.x == 0) {  
        array[idx] = 1;  
    }  
    __syncthreads();  
    array[idx] += threadIdx.x;  
}
```

If the kernel is launched with $\lll 1, 4 \ggg$ and array is initially {0, 0, 0, 0}, what is the value of array after execution?

- A. {1, 2, 3, 4}
- B. {1, 1, 1, 1}
- C. {1, 2, 3, 4}
- D. {1, 1, 2, 3}

Ans: D

Solution:

if(threadIdx.x == 0) condition is true only for thread 0. Thus array[0] is set to 1.

__syncthreads() function ensures that all threads wait until array[0] is updated and then all thread execute array[idx] += threadIdx.x.

threadIdx.x = 0: array[0] = 1+0 =1
threadIdx.x = 1: array[1] = 0+1 =1
threadIdx.x = 2: array[2] = 0+2 =2
threadIdx.x = 3: array[3] = 0+3 =3

Question 8:

A CUDA kernel with a 3D grid and block configuration is launched. Let total number of blocks be 72 and total number of threads be 576. Then

what can be the different ways to represent a block structure.

- A. (2, 2, 2)
- B. (3, 2, 3)
- C. (2, 4, 1)
- D. (1, 8, 0)

Ans: A, C

Solution:

Threads per block = $576 / 72 = 8$
so, we can see that $2*2*2=8$ and $1*8=8$

Question 9:

Consider a 2D CUDA kernel that computes on a 2D matrix M of size 2048x2048 with launch parameters $\langle\langle\langle(32,32),(64,64)\rangle\rangle\rangle$. Each CUDA thread operates only on one data point. On what data point will the thread with the threadIdx.x=0, threadIdx.y=0, blockIdx.x=10 and =blockIdx.y=10 operate on. Assume that thread ids along the x-dimension represent rows and thread ids along the y-dimension represent columns. Choose the correct option.

- A. M[10][10]
- B. M[32][32]
- C. M[640][640]
- D. None of the above

Ans: C

Solution:

Total number of threads in the grid is:

Thread per block X Blocks in grid = $(64 \times 64) \times (32 \times 32) = 2048 \times 2048$, This is equal to the size of the matrix M.

Global Row(x) = blockIdx.x X blockDim.x + threadIdx.x

Global Row(x) = $10 \times 64 + 0 = 640$

Global Column(y) = blockIdx.y X blockDim.y + threadIdx.y = $10 \times 64 + 0 = 640$

So thread operates on data point M[640][640]

Question 10:

Which among the following statement is FALSE:

- A. `gridDim.x,y,z` gives the number of blocks in a grid, in the x,y,z direction respectively.
- B. `blockDim.x,y,z` gives the number of threads in a block, in the x,y,z direction respectively
- C. `blockDim.x * gridDim.x` gives the number of threads in a grid in the x direction
- D. Each thread in a block has unique id given by system variable `gridIdx.x`

Ans: D

*******END*******