

Warp Scheduling and Divergence

Soumyajit Dey, Assistant Professor,
CSE, IIT Kharagpur

December 10, 2019



Course Organization

Topic	Week	Hours
Review of basic COA w.r.t. performance	1	2
Intro to GPU architectures	2	3
Intro to CUDA programming	3	2
Multi-dimensional data and synchronization	4	2
Warp Scheduling and Divergence	5	2
Memory Access Coalescing	6	2
Optimizing Reduction Kernels	7	3
Kernel Fusion, Thread and Block Coarsening	8	3
OpenCL - runtime system	9	3
OpenCL - heterogeneous computing	10	2
Efficient Neural Network Training/Inferencing	11-12	6



GPU can be viewed as an array of Streaming Multiprocessors (SMs) Each SM has the following elements

- ▶ Registers that can be partitioned among threads of execution
- ▶ Several Caches: Shared memory, Constant, Texture, L1 etc
- ▶ Warp Schedulers (More on this later)
- ▶ Scalar Processors (SPs) for integer and floating-point operations
- ▶ Special Function Units (SFUs) for single-precision floating-point transcendental functions

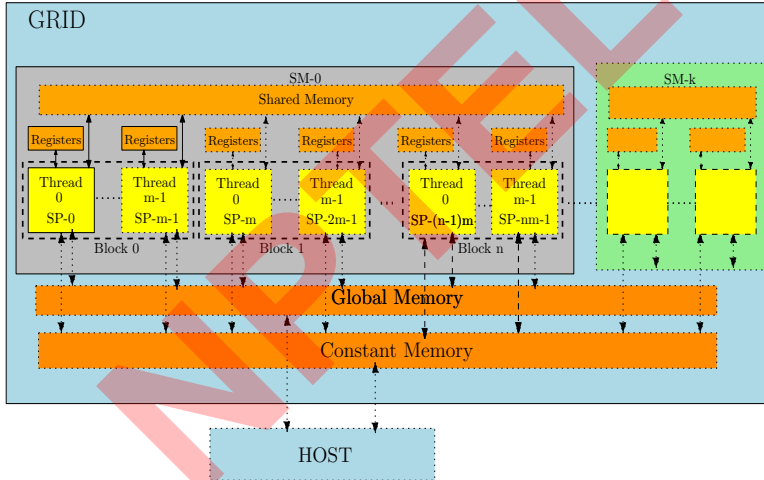


Table: CUDA Device Memory Types and Scopes

Variables Declaration	Memory	Scope	Lifetime
Automatic Variables other than arrays	Register	Thread	Kernel
Automatic array variables	Local	Thread	Kernel
<code>__device__ __shared__ int SharedVar</code>	Shared	Block	Kernel
<code>__device__ int GlobalVar</code>	Global	Grid	Application
<code>__device__ __constant__ int ConstVar</code>	Constant	Grid	Application



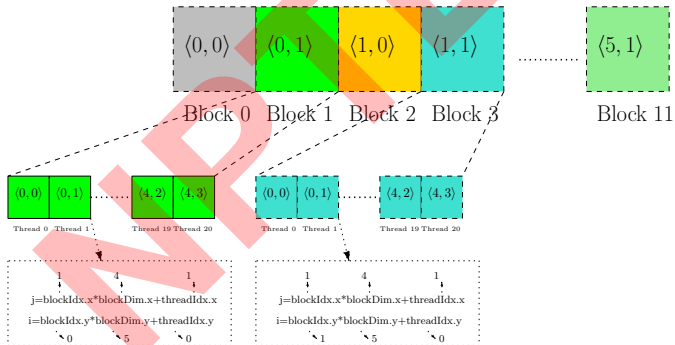
Mapping to Hardware



Example: CUDA Thread and Block Definition

$$\text{NumCols} = \text{blockDim.x} * \text{gridDim.x}$$

```
NumRows = blockDim.y * gridDim.y
```

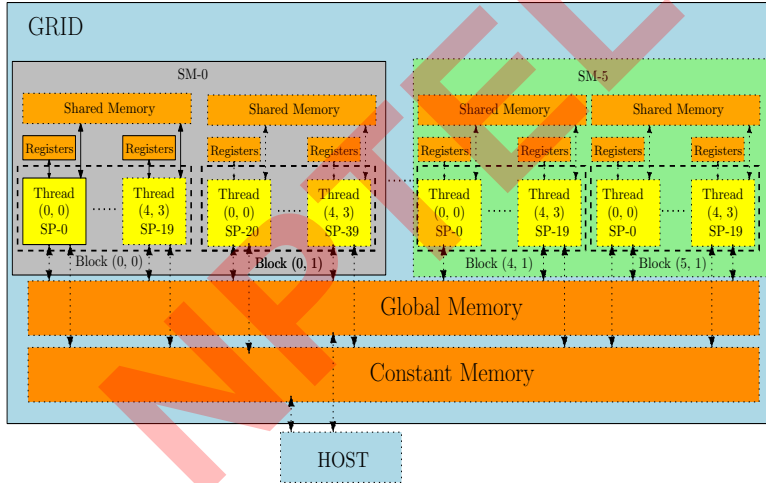
$$gridDim = \langle 6, 3 \rangle$$
$$blockDim = \langle 4, 5 \rangle$$


Generalized Mapping Scenario

- ▶ Let us consider a scenario for the grid and block dimensions specified above.
- ▶ $gridDim = \langle 6, 2 \rangle$ and $blockDim = \langle 5, 4 \rangle$
- ▶ $\#SMs = 6$ $\#SPs$ per $SM = 40$
- ▶ Two Blocks are mapped to one SM at a time.
- ▶ Hardware resources are completely utilized.



Mapping to Hardware

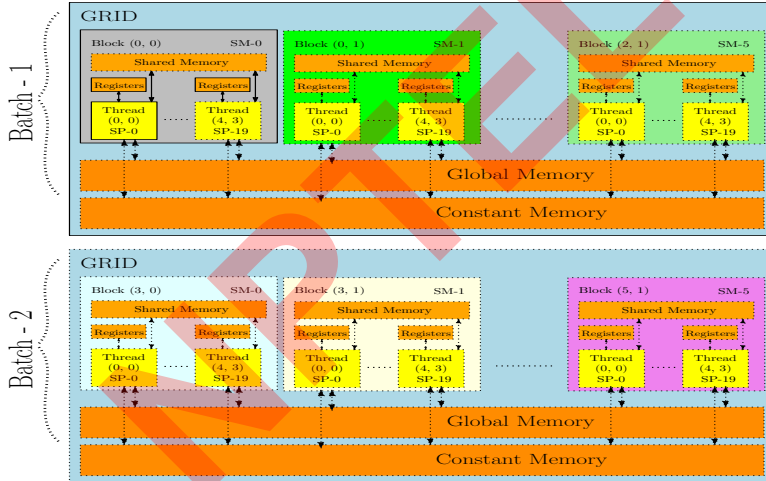


Mapping in a resource constrained setting

- ▶ Consider a scenario where the resources of the architecture are limited.
- ▶ $gridDim = \langle 6, 2 \rangle$ and $blockDim = \langle 5, 4 \rangle$
- ▶ $\#SMs = 6$ $\#SPs$ per $SM = 20$
- ▶ Thread Blocks are launched in batches sequentially.
- ▶ Execution is serialized to some extent.



Mapping to Hardware



SM, SP, Block and thread

- ▶ thread block max size : 1024 (modern archs 2048)
- ▶ SM can store max 1024 "thread contexts"
- ▶ can have much less than 1024 SPs
- ▶ GTX 970 : 13 SMs : 13 X 1024 thread contexts in parallel
- ▶ GTX 970 : 128 SP per SM



SM, SP, Block and thread

- ▶ One block in one SM
- ▶ One SM can have multiple blocks

If SM can store max 1024 "thread contexts", and block size is 256, we have 4 blocks per SM.

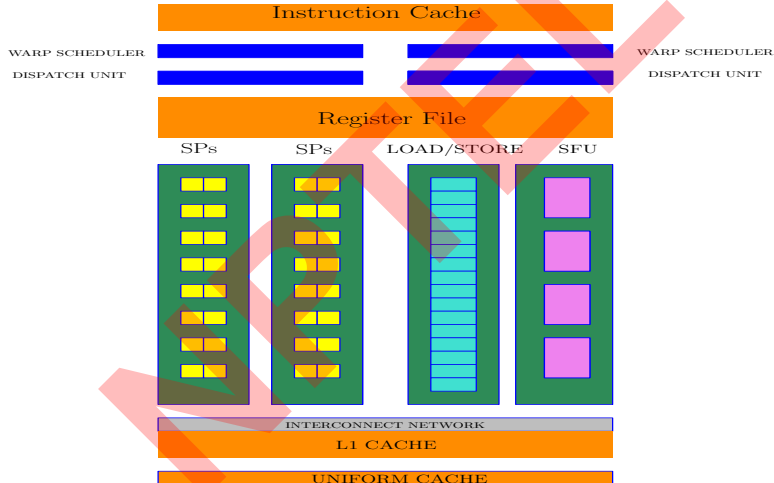


GPU HW scheduler

- ▶ The hw scheduler decided which threads to map to a collection of SPs in SIMD fashion :: SIMT model of execution
- ▶ This collection is physically guaranteed to execute in parallel
- ▶ The unit of such collections is "warp"



SM: A closer look



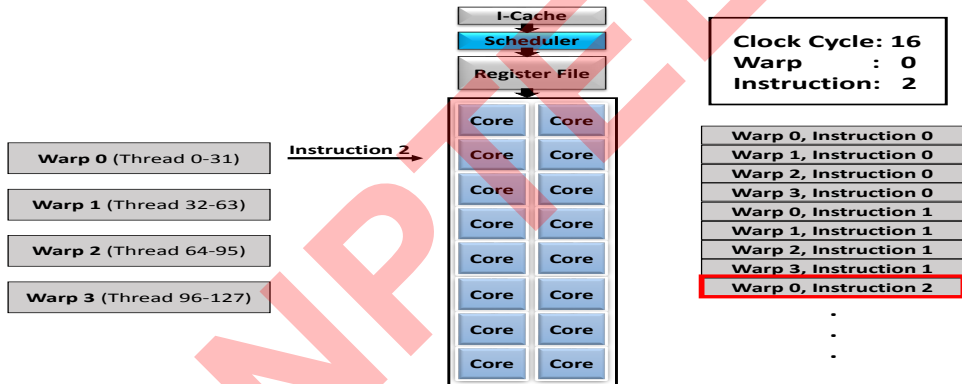
Warps

- ▶ Warp is a unit of thread Scheduling in SMs.
- ▶ Warp size is implementation specific (typically 32 threads)
- ▶ Warps are executed in an SIMD fashion i.e. the warp scheduler launches warps of threads and each warp typically executes one instruction across parallel threads.

Ex : If a SM has 128 SPs, it can execute 4 Warps at a given time (one Warp has 32 Threads)



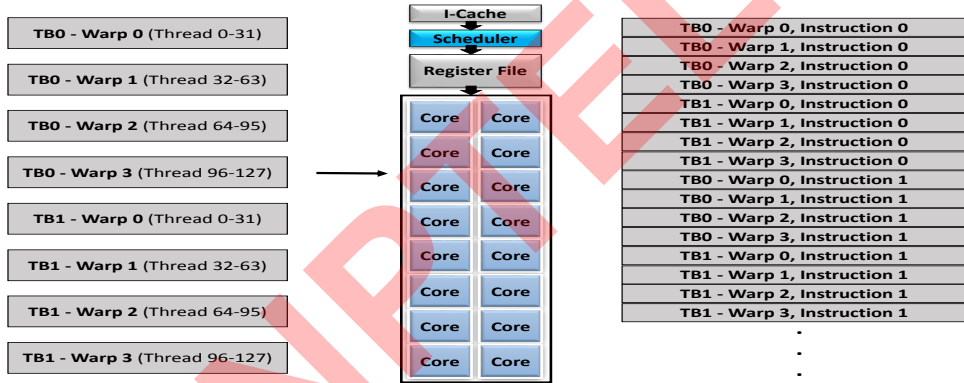
Warp Scheduling in SM



Ref : Henk Corporaal, Gert-Jan van den Braak - "Introduction to GPGPU Architectures"




Warp Scheduling in SM



- ▶ Thread block scheduler (TBS) is *believed to use* round robin policy to schedule thread blocks - implementation dependent



Warp Scheduling in SM



```
__global__ void twice(int *array)
{
    int tid = blockDim.x*blockIdx.x+threadIdx.x;
    array[tid] = array[tid] + array[tid];
}
```

```
mov.u32    %r1, %ntid.x;
mov.u32    %r2, %ctaid.x;
mov.u32    %r3, %tid.x;
mad.lo.s32 %r4, %r2, %r1, %r3;
mul.wide.s32 %rd3, %r4, 4;
add.s64    %rd4, %rd2, %rd3;
ld.global.u32 %r5, [%rd4];
add.s32    %r6, %r5, %r5;
st.global.u32 [%rd4], %r6;
```

Figure: Simple CUDA Kernel



Warp Scheduling in SM

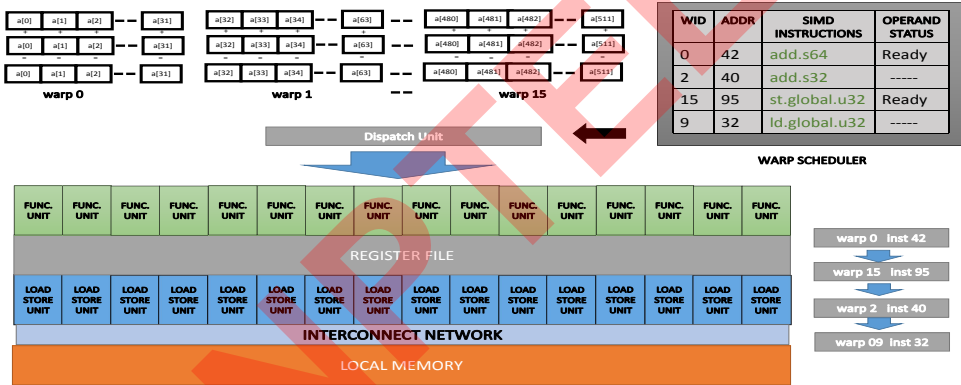


Figure: Warp Scheduler



Warp Scheduling in SM

- ▶ Issue one “ready-to-go” warp instruction/cycle
- ▶ Use operand score-boarding to prevent hazards
- ▶ Issue selection based on round-robin/age of warp
- ▶ Score-boarding determines if a thread is ready to execute?
- ▶ Scoreboard is a HW implemented table that tracks - instrs fetched, resource availability for fetched instrs (FU and operand), register file modifications by instrs.



Latency Tolerance

- ▶ When threads in one warp execute a long-latency operation (read from global memory), the warp scheduler will dispatch and execute other warps until that operation is finished.
- ▶ Other long latency operations : FP units, Branch instructions
- ▶ After all, all threads in the same control-flow execute same instruction sequence on different data points !
- ▶ A common practice is to launch thread blocks of a size that is a multiple of the warp size to maximally utilize threads.
- ▶ Slow global memory accesses by threads in a warp may be optimized using coalescing (more on this later)



Efficient use of thread blocks

Target System Constraints

- ▶ A maximum of 8 blocks and 1024 threads per SM
- ▶ A maximum of 512 threads per block

Table: Solutions for various block scenarios

Input Block Size	Blocks per SM	Threads per Block	Remarks
8 * 8	12	64	SM execution resources will be underutilized
16*16	4	256	Achieves full thread capacity in SMs
32*32	1	1024	Exceeds the limit of 512 threads per block



Querying Device Properties

CUDA API provides constructs for obtaining properties of the target GPU.

- ▶ **cudaGetDeviceCount()**: Obtains the number of devices in the system.
- ▶ **cudaGetDeviceProperties()**: Returns the property values of a particular device



Querying Device Properties

```
int main()
{
    int devCount;
    cudaGetDeviceCount(&devCount);
    for (int i = 0; i < devCount; ++i)
    {
        cudaDeviceProp devp;
        cudaGetDeviceProperties(&devp, i);
        printDevProp(devp);
    }
    return 0;
}
```



Querying Device Properties

```
void printDevProp(cudaDeviceProp devProp)
{
    printf("Major revision number: %d\n",devProp.major);
    printf("Minor revision number: %d\n",devProp.minor);
    printf("Name: %s\n",devProp.name);
    printf("Total global memory: u\n",devProp.totalGlobalMem);
    printf("Total shared memory per block:%u\n", devProp.sharedMemPerBlock);
    printf("Total registers per block: %d\n", devProp.regsPerBlock);
    printf("Warp size: %d\n",devProp.warpSize);
    printf("Maximum memory pitch: %u\n",devProp.memPitch);
    printf("Maximum threads per block: %d\n",devProp.maxThreadsPerBlock);
    for (int i = 0; i < 3; ++i)
        printf("Maximum dimension %d of block: %d\n",i,devProp.maxThreadsDim[i]);
    for (int i = 0; i < 3; ++i)
        printf("Maximum dimension %d of grid: %d\n", i, devProp.maxGridSize[i]);
}
```



Querying Device Properties

```
printf("Clock rate: %d\n",devProp.clockRate);  
printf("Total constant memory:%u\n", devProp.totalConstMem);  
printf("Texture alignment: %u\n", devProp.textureAlignment);  
printf("Concurrent copy and execution: %s\n", (devProp.deviceOverlap ? "Yes"  
    : "No"));  
printf("Number of multiprocessors: %d\n",devProp.multiProcessorCount);  
return;  
}
```



Example: Tesla K40m Characteristics

Major revision number: 3
Minor revision number: 5
Name: Tesla K40m
Total global memory: 3405643776
Total shared memory per block: 49152
Total registers per block: 65536
Warp size: 32
Maximum memory pitch: 2147483647
Maximum threads per block: 1024
Maximum dimension 0 of block: 1024
Maximum dimension 1 of block: 1024
Maximum dimension 2 of block: 64
Maximum dimension 0 of grid: 2147483647
Maximum dimension 1 of grid: 65535
Maximum dimension 2 of grid: 65535
Clock rate: 745000
Total constant memory: 65536
Texture alignment: 512
Concurrent copy and execution: Yes
Number of multiprocessors: 15



Control Flow Divergence

- ▶ Threads inside a warp execute the same instruction.
- ▶ How does a warp handle if statements / branch instructions?
- ▶ The GPU is not capable of running both the if else blocks at the same time.



Warp Scheduling

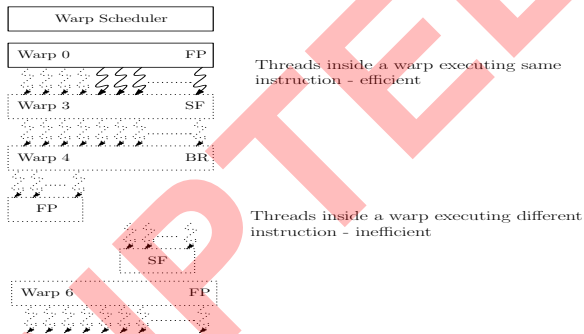


Figure: Warp Divergence



Divergent Code 1

Consider the following kernel code

```
__global__  
void divergence(float *M)  
{  
    /*P1:*/ int tid=blockIdx.x*blockDim.x+threadIdx.x;  
    /*P2:*/ if(tid%2)  
    /*P3:*/     M[j]+=2;  
    else  
    /*P4:*/     M[j]-=2;  
    /*P5:*/ M[j]*=2;  
}
```

Half the threads of a warp execute the addition instruction while the other half execute the subtraction instruction.




The Hardware's Job

The GPU has hardware support for handling divergent branch instructions in code.

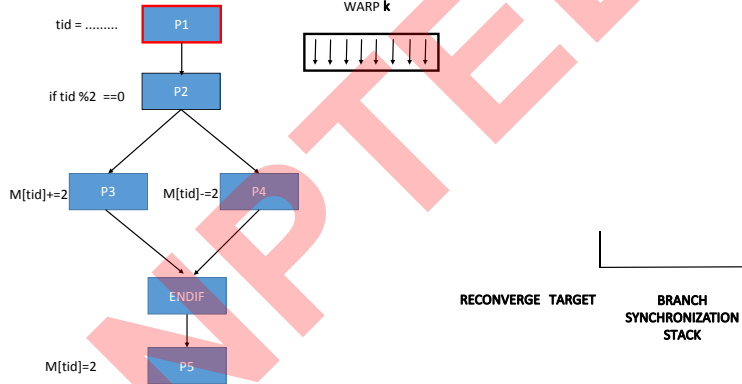
- ▶ The PTX Assembler maintains internal masks, a branch synchronization stack and special markers
- ▶ The PTX Assembler sets a **branch synchronization marker** first for the divergent `if` statement that pushes the active mask on a stack inside each SIMD thread
- ▶ Depending on the value of the mask relevant threads execute instructions,
- ▶ Once the instructions in the `if` block are finished, the active mask is popped from the stack, flipped and pushed back.



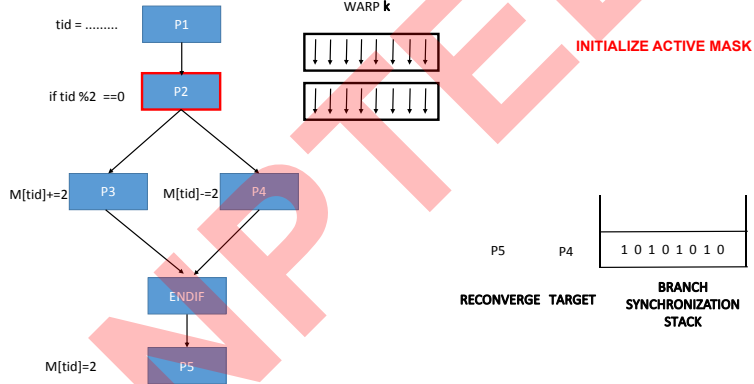
Divergent Code 1



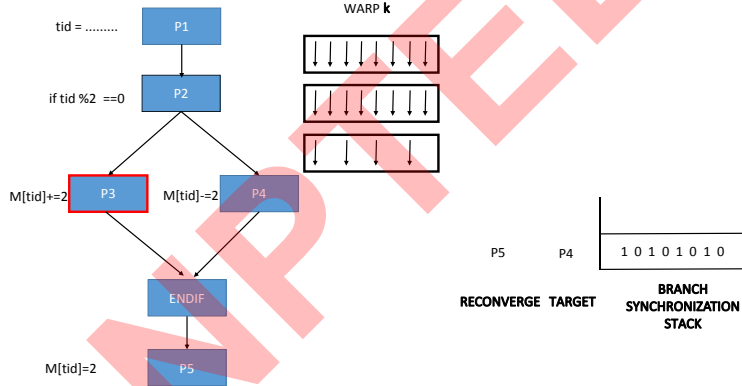
Divergent Code 1



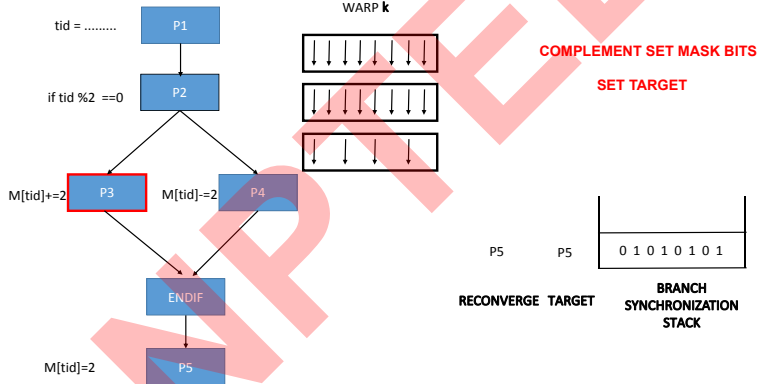
Divergent Code 1



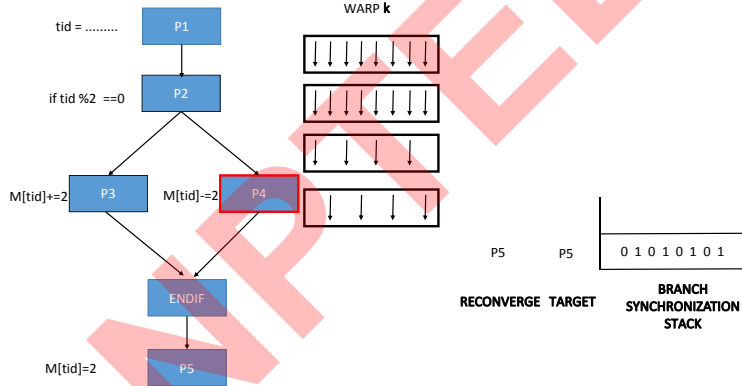
Divergent Code 1



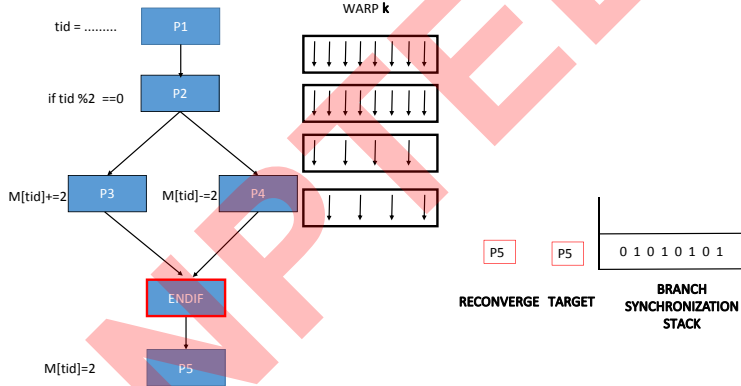
Divergent Code 1



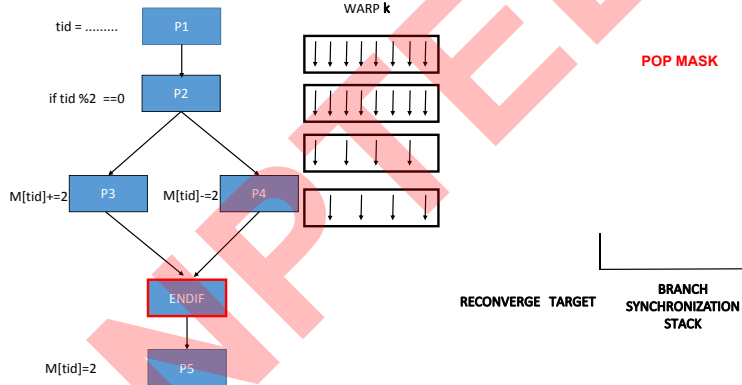
Divergent Code 1



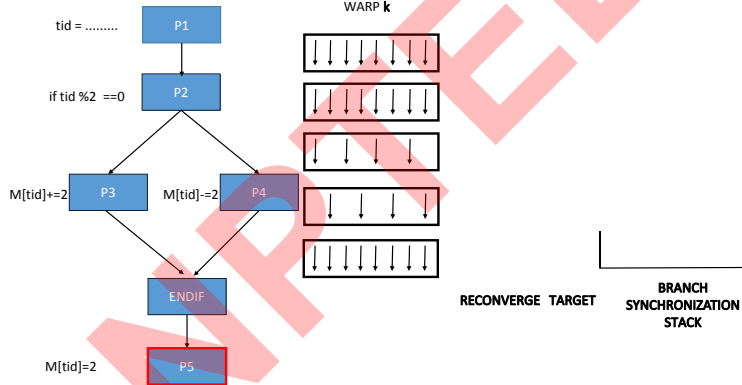
Divergent Code 1



Divergent Code 1



Divergent Code 1




Divergent Code 2

Let us consider an example that has nested if/else statements.

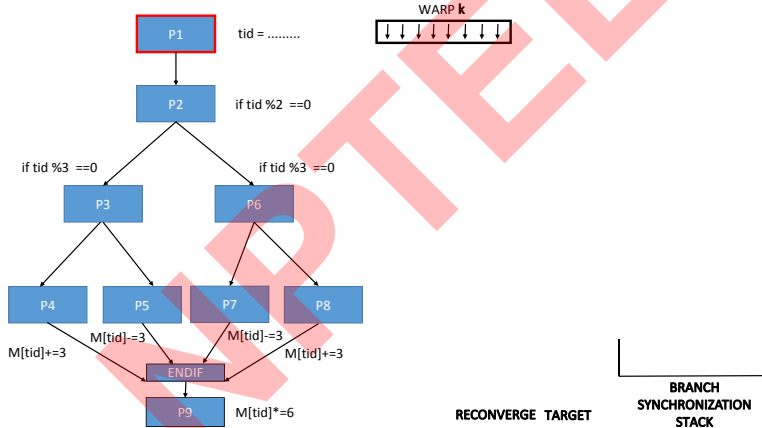
```
__global__
void divergence(float *M)
{
    /*P1*/    int tid=blockIdx.x*blockDim.x+threadIdx.x;
    /*P2*/    if(tid%2==0)
    {
        /*P3*/    if(tid%3==0)
        /*P4*/        M[tid]+=3;
        else
        /*P5*/        M[tid]-=3;
    }
    else
    {
        /*P6*/    if(tid%3==0)
        /*P7*/        M[tid]-=3;
        else
        /*P8*/        M[tid]+=3;
    }
    /*P9*/    M[tid]*=6;
}
```



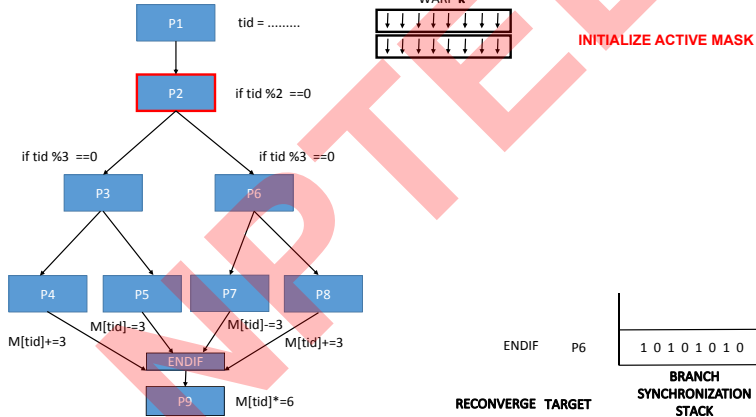
Divergence Code 2





Divergence Code 2




Divergence Code 2



Divergence Code 2





ENDIF P6



RECONVERGE TARGET



Divergence Code 2





PUSH OLD MASK

ENDIF	P5	1 0 1 0 1 0 1 0
P9	P6	1 0 1 0 1 0 1 0

RECONVERGE TARGET



Divergence Code 2



SET RELEVANT BITS

ENDIF P5


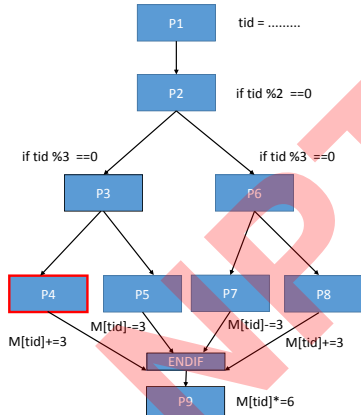
ENDIF P6



RECONVERGE TARGET



Divergence Code 2





ENDIF	P5	0	0	0	0	0	1	0
ENDIF	P6	1	0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2





ENDIF	P5	0	0	0	0	0	1	0
ENDIF	P6	1	0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2



COMPLEMENT
RELEVANT BITS



ENDIF P5
ENDIF P6

0	0	0	0	0	0	1	0
1	0	1	0	1	0	1	0

RECONVERGE TARGET



Divergence Code 2



COMPLEMENT
RELEVANT BITS



ENDIF ENDIF
ENDIF P6

1	0	1	0	1	0	0	0
1	0	1	0	1	0	1	0

RECONVERGE TARGET



Divergence Code 2



COMPLEMENT
RELEVANT BITS



ENDIF ENDIF
ENDIF P6

1	0	1	0	1	0	0	0
1	0	1	0	1	0	1	0

RECONVERGE TARGET




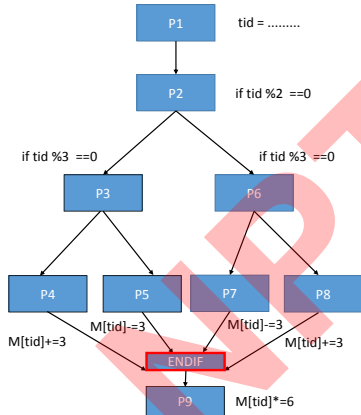
Divergence Code 2



RECONVERGE TARGET



Divergence Code 2



POP STACK

ENDIF



P6

1 0 1 0 1 0 1 0

RECONVERGE TARGET



Divergence Code 2



ENDIF



P6

1 0 1 0 1 0 1 0

RECONVERGE TARGET



Divergence Code 2



COMPLEMENT
RELEVANT BITS
SET TARGET



ENDIF ENDF

0 1 0 1 0 1 0 1

RECONVERGE TARGET



Divergence Code 2





ENDIF ENDIF

0 1 0 1 0 1 0 1

RECONVERGE TARGET



Divergence Code 2





PUSH OLD MASK

ENDIF	P8	0 1 0 1 0 1 0 1
ENDIF	ENDIF	0 1 0 1 0 1 0 1

RECONVERGE TARGET



Divergence Code 2





ENDIF P8
ENDIF ENDIF

0	0	1	0	0	0
0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2




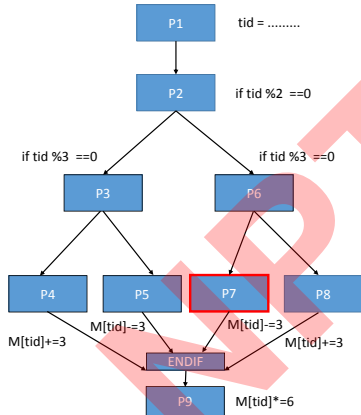
ENDIF P8
ENDIF ENDIF

0	0	1	0	0	0
0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2



COMPLEMENT
RELEVANT BITS

SET TARGET



ENDIF ENDIF
ENDIF ENDIF

0	1	0	0	0	1	0	1
0	1	0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2





ENDIF	ENDIF	0	1	0	0	0	1	0	1
ENDIF	ENDIF	0	1	0	1	0	1	0	1

RECONVERGE TARGET



Divergence Code 2


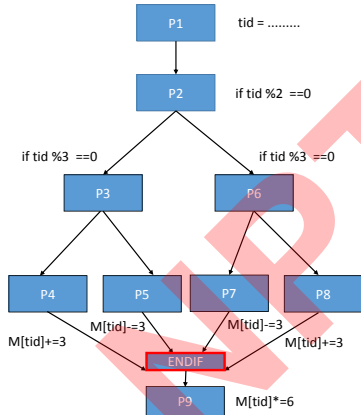


ENDIF	ENDIF	0	1	0	0	0	1	0	1
ENDIF	ENDIF	0	1	0	1	0	1	0	1


RECONVERGE TARGET



Divergence Code 2




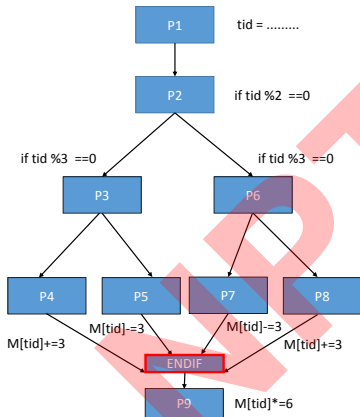
POP STACK



RECONVERGE TARGET



Divergence Code 2




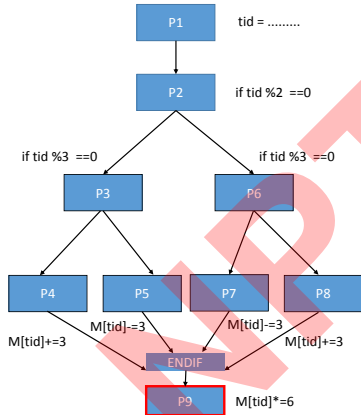
ENDIF ENDIF

0 1 0 1 0 1 0 1

RECONVERGE TARGET



Divergence Code 2



Programming tips

- ▶ GPU programmer has to be aware of hardware imposed restrictions - threads/SM, blocks/SM, threads/blocks, threads/warps
- ▶ The only safe way to synchronize threads from different blocks is to terminate kernel and make a fresh launch at the target synchronization point

